

Two very different firms shouldn't score almost the same. Right now, they do.

The matching algorithm's quality score is built entirely from our internal compliance Assessment — bar standing, disciplinary history, engagement process. That measures whether a firm is legitimate, not whether it's actually good. This proposal closes that gap with one independently-verifiable signal per firm, and a concrete plan for how it moves real rankings.



THE ASK

Okay to have Liam and Olivia each look up one external number per firm — Google rating and review count — for every firm in the database?

The problem, with real numbers

We already hit this exact issue once this cycle and fixed it — worth showing, because it's the same fix this proposal extends. For a real pool of New York contract-law firms, the quality score used to spread firms across only 12.2–17.9 of a possible 28 points — a 5.7-point band standing in for what should be a full range of firm quality. The best and the fifth-best firm in the same search barely separated.

Reworking quality to score firms *relative to the actual pool being compared*, instead of on a flat scale, immediately fixed it — same firms, same data, now spread the full 0.0–28.0 range. Real quality differences the algorithm already had access to just weren't visible before.

That fix helps whatever data feeds it separate firms further. It can't invent differentiation that isn't there. Chambers/Super Lawyers/AV data — extracted below — only exists for 43% of firms. A rating every firm has is what makes the fix pay off everywhere, not just where a firm happens to already have a ranking on file.

What's already covered, no new research needed

Before assigning anything, I ran an extraction pass against the firm data we already have — a lot of firms already reference Chambers rankings, Super Lawyers, Best Lawyers, or AV Preeminent status in their existing write-ups, just as unstructured text. Pulling that into structured fields required no manual lookup at all — it populates five more columns for free.

170

of 392 firms (43%) get a bonus signal layered on top — Chambers/Super Lawyers/AV data, extracted automatically

100

firms have a Chambers ranking mention on file

392

firms get the Google rating lookup — every firm, for a consistent baseline signal

392 firms

■ Also has Chambers/SL/AV data (170) ■ Google rating is their only external signal (222)

Why every firm gets the Google lookup rather than just the 222 gaps: mixed signal availability would distort the blend — a firm scored on two external signals shouldn't be structurally different from one scored on none just because of what we happened to already have on file. One consistent baseline, everywhere.

What's actually left for Liam & Olivia

One field, one source, every firm, split by batch so there's no duplicated work.

Liam

First half of the firm list · search "[firm name] [city] reviews," record the Google rating + review count

Olivia

Second half of the firm list · same task

google_rating

google_review_count

Two new columns on the existing firm sheet. No new tooling, no new process to learn.

How this actually changes rankings, not just sits as a column

This plugs into `scoreQuality` in the matching algorithm directly, not a separate report:

- ✓ **Review count weights the rating.** A 5.0 from 3 reviews won't outrank a 4.6 from 200 — low-review firms get pulled toward the dataset average in proportion to how little signal they actually have, so a handful of reviews can't distort the ranking.
- ✓ **Blended with the Chambers/Super Lawyers/AV data** as a secondary boost, for the 170 firms that have it.
- ✓ **Run through the pool-relative scoring already shipped this cycle** — quality is scored relative to the other firms in each specific search, not on a flat 0–100 scale, so real variance in this new data visibly separates firms instead of bunching everyone near the top.

Keeping the research itself simple

- ✓ One field per firm, not a multi-source research task — no Chambers/Avvo directory digging (tried it firsthand; too many firms aren't reliably searchable there).
- ✓ Blank is an acceptable answer. If a firm has no Google listing, they move on — no chasing a number that doesn't exist.
- ✓ Split by firm batch, not by task, so there's zero coordination overhead between the two of them.